

Univariate Linear Regression

Load the Data and Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns

import matplotlib.pyplot as plt
plt.style.use('ggplot')
%matplotlib inline
```

```
In [2]: df_bike = pd.read_csv("bike_sharing_data.txt")
df_bike.head()
```

Out[2]:

	Population	Profit
0	6.1101	17.5920
1	5.5277	9.1302
2	8.5186	13.6620
3	7.0032	11.8540
4	5.8598	6.8233


```
In [3]: df_bike.shape
```

Out[3]: (97, 2)

Visualize the Data

```
In [4]: fig = plt.figure(figsize = (20, 8))
ax = sns.scatterplot(x = "Population", y = "Profit", data = df_bike)
ax.set_title("Profit in $10000s vs City Population in 10000s")
```

```
Out[4]: <matplotlib.text.Text at 0x90961f0>
```



Compute the Cost

The objective of linear regression is to minimize the cost function

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

where $h_{\theta}(x)$ is the hypothesis and given by the linear model

$$h_{\theta}(x) = \theta^T x = \theta_0 + \theta_1 x_1$$

```
In [5]: def cost_function(X, y, theta):
        m = len(y)
        y_pred = X.dot(theta)
        error = (y_pred - y) ** 2

        return 1 / (2 * m) * np.sum(error)
```

```
In [6]: m = df_bike.Population.values.size
X = np.append(np.ones((m, 1)), df_bike.Population.values.reshape(m, 1), axis = 1)
y = df_bike.Profit.values.reshape(m, 1)
theta = np.zeros((2, 1))

cost_function(X, y, theta)
```

Out[6]: 32.072733877455676

Gradient Descent

Minimize the cost function $J(\theta)$ by updating the below equation and repeat until convergence

$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$ (simultaneously update θ_j for all j).

```
In [7]: def gradient_descent(X, y, theta, alpha, iterations):
        m = len(y)
        costs = []
        for i in range(iterations):
            y_pred = X.dot(theta)
            error = np.dot(X.transpose(), (y_pred - y))
            theta -= alpha * 1/m * error
            costs.append(cost_function(X, y, theta))
        return theta, costs
```

```
In [8]: theta, costs= gradient_descent(X, y, theta, alpha = 0.01, iterations = 2000)

print(" h(x) = {} + {}x1 ".format(str(round(theta[0, 0], 2)),
                                   str(round(theta[1, 0], 2))))

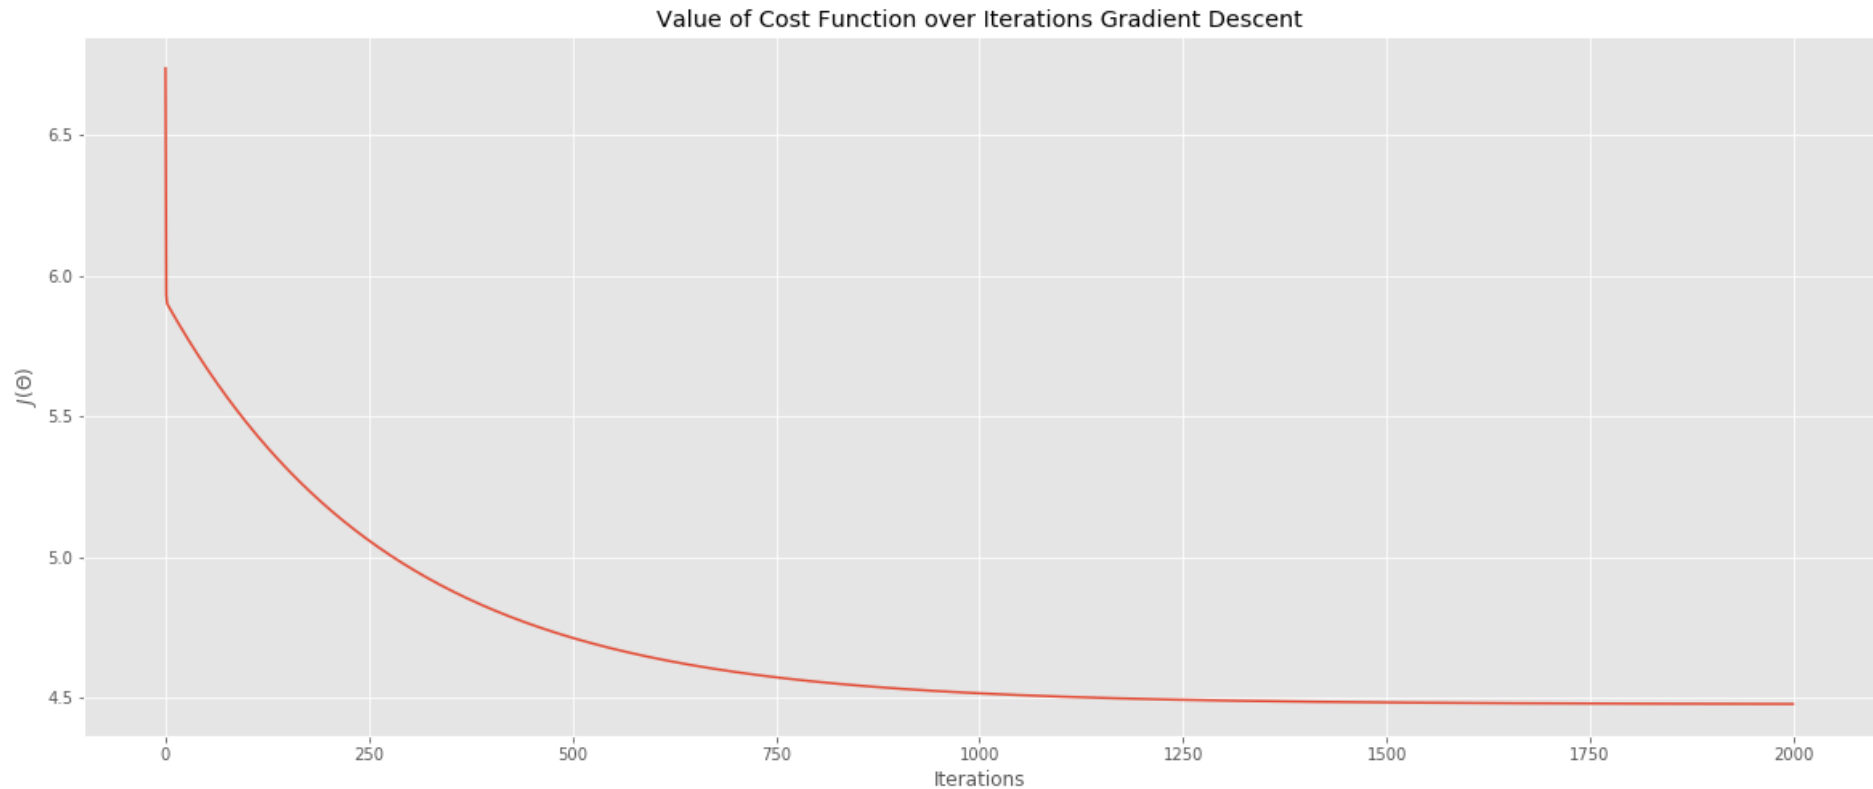
h(x) = -3.79 + 1.18x1
```

Plotting the Convergence

Plot $J(\theta)$ against the number of iterations of gradient descent:

```
In [9]: fig = plt.figure(figsize = (20, 8))  
plt.plot(costs)  
plt.xlabel("Iterations")  
plt.ylabel("$J(\theta)$")  
plt.title("Value of Cost Function over Iterations Gradient Descent")
```

Out[9]: <matplotlib.text.Text at 0x94890b0>



Training Data with Linear Regression Fit

```
In [10]: theta.shape
```

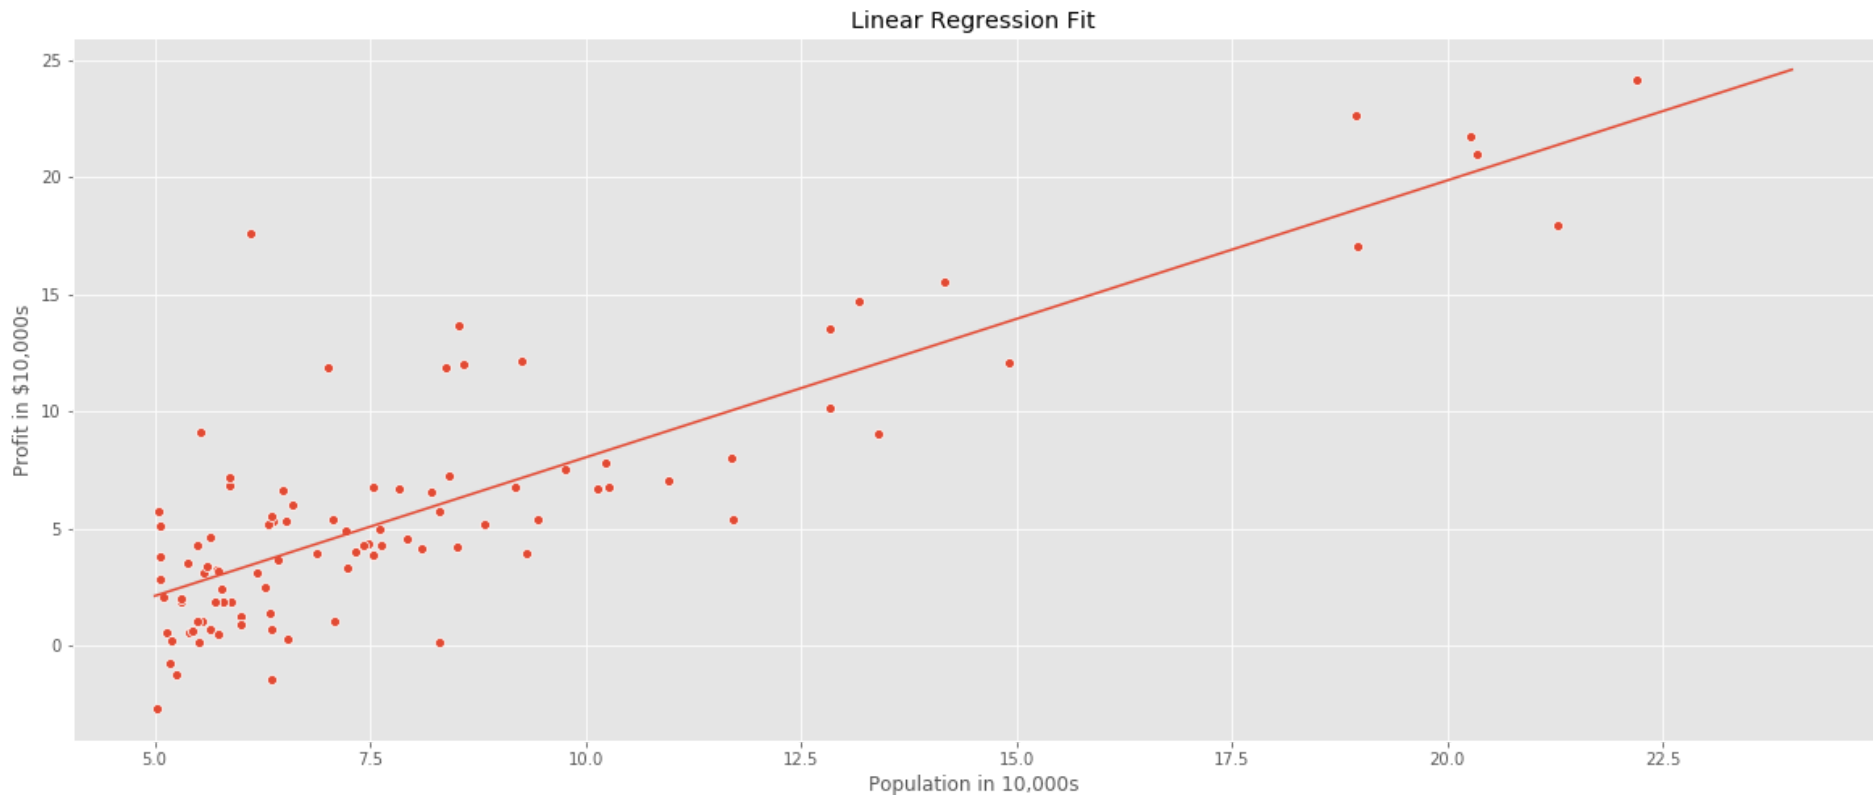
Out[10]: (2, 1)

```
In [11]: theta
```

```
Out[11]: array([[ -3.78806857],  
               [ 1.18221277]])
```

```
In [12]: fig = plt.figure(figsize = (20, 8))  
theta = np.squeeze(theta)  
ax = sns.scatterplot(x = "Population", y = "Profit", data = df_bike)  
  
x_value = [X for X in range(5, 25)]  
y_value = [(X * theta[1] + theta[0]) for X in x_value]  
sns.lineplot(x_value, y_value)  
  
plt.xlabel("Population in 10,000s")  
plt.ylabel("Profit in $10,000s")  
plt.title("Linear Regression Fit")
```

```
Out[12]: <matplotlib.text.Text at 0x9408110>
```



Inference using the optimized θ values

$$h_{\theta}(x) = \theta^T x$$

```
In [13]: def predict(X, theta):  
        y_pred = np.dot(theta.transpose(), X)  
        return y_pred
```

```
In [14]: y_pred_1 = predict(np.array([1, 4]), theta) * 10000  
print("For a population of 40,000 people, the model predicts a profit of $" + str(round(y_pred_1, 0)))
```

For a population of 40,000 people, the model predicts a profit of \$9408.0

```
In [15]: y_pred_2 = predict(np.array([1, 8.3]), theta) * 10000  
print("For a population of 83,000 people, the model predicts a profit of $" + str(round(y_pred_2, 0)))
```

For a population of 83,000 people, the model predicts a profit of \$60243.0